

1. Select the option that correctly completes the sentence: 1 point

Training a model using labeled data and using this model to predict the labels for new data is known as \_\_\_\_\_.

- Unsupervised Learning
- Clustering
- Supervised Learning
- Density Estimation

2. Select the option that correctly completes the sentence: 1 point

Modeling the features of an unlabeled dataset to find hidden structure is known as \_\_\_\_\_.

- Unsupervised Learning
- Classification
- Supervised Learning
- Regression

3. Select the option that correctly completes the sentence: 1 point

Training a model using categorically labelled data to predict labels for new data is known as \_\_\_\_\_.

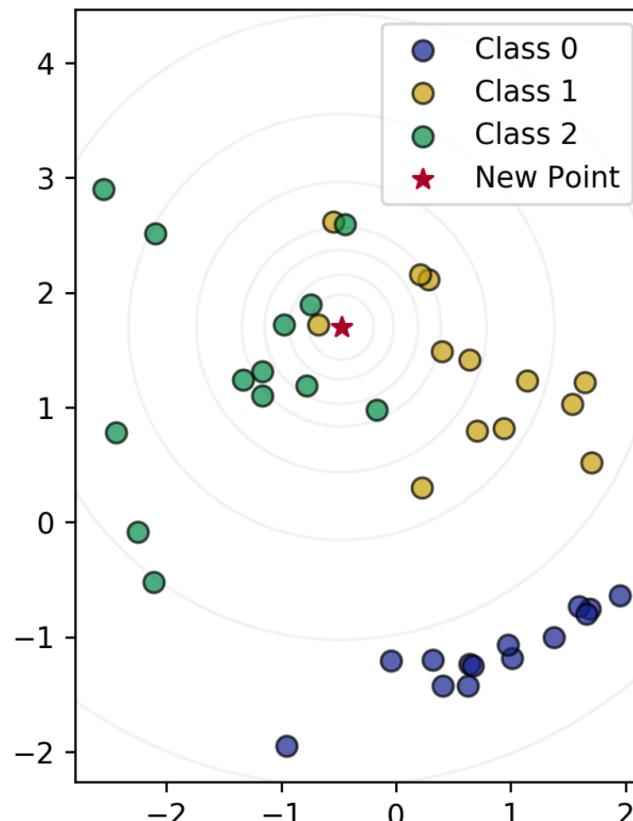
- Clustering
- Feature Extraction
- Classification
- Regression

4. Select the option that correctly completes the sentence: 1 point

Training a model using labelled data where the labels are continuous quantities to predict labels for new data is known as \_\_\_\_\_.

- Clustering
- Feature Extraction
- Regression
- Classification

5. Using the data for classes 0, 1, and 2 plotted below, what class would a KNeighborsClassifier classify the new point as for k = 1 and k = 3? 1 point



- k=1: Class 0
- k=3: Class 1
- k=1: Class 2
- k=3: Class 1
- k=1: Class 1
- k=3: Class 2
- k=1: Class 0
- k=3: Class 2
- k=1: Class 1
- k=3: Class 0

6. Which of the following would \*not\* be something you specify for a nearest neighbor classifier algorithm?

1 point

- A distance metric that finds neighbor points in the training set
- How many nearest neighbors to examine
- A method for pooling the classes of neighbor points to make a final classification decision
- An optional weighting function on the neighbor points
- The number of initial clusters to create from the training set

7. Why is it important to examine your dataset as a first step in applying machine learning? (Select all that apply):

1 point

- See what type of cleaning or preprocessing still needs to be done
- You might notice missing data
- Gain insight on what machine learning model might be appropriate, if any
- Get a sense for how difficult the problem might be
- It is not important

8. The key purpose of splitting the dataset into training and test sets is:

1 point

- To reduce the amount of labelled data needed for evaluating classifier accuracy
- To reduce the number of features we need to consider as input to the learning algorithm
- To speed up the training process
- To estimate how well the learned model will generalize to new data

9. The purpose of setting the random\_state parameter in train\_test\_split is: (Select all that apply)

1 point

- To split the data into similar subsets so that bias is not introduced into the final results
- To make experiments easily reproducible by always using the same partitioning of the data
- To avoid bias in data splitting
- To avoid predictable splitting of the data

10. Given a dataset with 10,000 observations and 50 features plus one label, what would be the dimensions of X\_train, y\_train, X\_test, and y\_test? Assume a train/test split of 75%/25%.

1 point

- X\_train: (7500, 50)  
• y\_train: (7500, )  
• X\_test: (2500, 50)  
• y\_test: (2500, )
- X\_train: (2500, )  
• y\_train: (2500, 50)  
• X\_test: (7500, )  
• y\_test: (7500, 50)
- X\_train: (10000, 28)  
• y\_train: (10000, )  
• X\_test: (10000, 12)  
• y\_test: (10000, )
- X\_train: (2500, 50)  
• y\_train: (2500, )  
• X\_test: (7500, 50)  
• y\_test: (7500, )

- X\_train: (10000, 50)  
• y\_train: (10000, )  
• X\_test: (10000, 50)  
• y\_test: (10000, )

11. Which of the following statements about the effect of the number of neighbors parameter k is true?

1 point

- Nearest neighbor classifier accuracy is never significantly affected by the choice of k.  
 For small values of k (e.g. k = 1), the classifier will tend to be much more sensitive to noise, mislabeled data, and other sources of variation for individual data points.  
 For larger values of k (e.g. k = 10), the regions in feature space assigned to different classes generally become less fragmented and have smoother decision boundaries between the regions, compared to results using smaller values of k.  
 For small values of k (e.g. k = 1), the classifier will tend to be much less sensitive to noise, mislabeled data, and other sources of variation for individual data points.